

```
In [4]:  import pandas as pd
import numpy as np
```

```
In [6]:  df = pd.read_csv('Housing.csv')
```

```
In [7]:  df
```

Out[7]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hot
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	
...
540	1820000	3000	2	1	1	yes	no	yes	
541	1767150	2400	3	1	1	no	no	no	
542	1750000	3620	2	1	1	yes	no	no	
543	1750000	2910	3	1	1	no	no	no	
544	1750000	3850	3	1	2	yes	no	no	

545 rows × 13 columns



```
In [4]:  ## DATA PREPROCESSING ##
```

```
In [8]:  df.head()
```

Out[8]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwa
0	13300000	7420	4	2	3	yes	no	no	
1	12250000	8960	4	4	4	yes	no	no	
2	12250000	9960	3	2	2	yes	no	yes	
3	12215000	7500	4	2	2	yes	no	yes	
4	11410000	7420	4	1	2	yes	yes	yes	



In [9]: `df.tail()`

Out[9]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotw
540	1820000	3000	2	1	1	yes	no	yes	
541	1767150	2400	3	1	1	no	no	no	
542	1750000	3620	2	1	1	yes	no	no	
543	1750000	2910	3	1	1	no	no	no	
544	1750000	3850	3	1	2	yes	no	no	

In [10]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 545 entries, 0 to 544
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                 545 non-null    int64
1   area                  545 non-null    int64
2   bedrooms              545 non-null    int64
3   bathrooms             545 non-null    int64
4   stories               545 non-null    int64
5   mainroad              545 non-null    object
6   guestroom             545 non-null    object
7   basement              545 non-null    object
8   hotwaterheating       545 non-null    object
9   airconditioning       545 non-null    object
10  parking               545 non-null    int64
11  prefarea              545 non-null    object
12  furnishingstatus      545 non-null    object
dtypes: int64(6), object(7)
memory usage: 55.5+ KB
```

In [11]: `df.shape`

Out[11]: (545, 13)

In [12]: `df.columns`

Out[12]: Index(['price', 'area', 'bedrooms', 'bathrooms', 'stories', 'mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', 'parking', 'prefarea', 'furnishingstatus'], dtype='object')

In [13]: `df.duplicated().sum()`

Out[13]: 0

```
In [14]: df.isnull().sum()
```

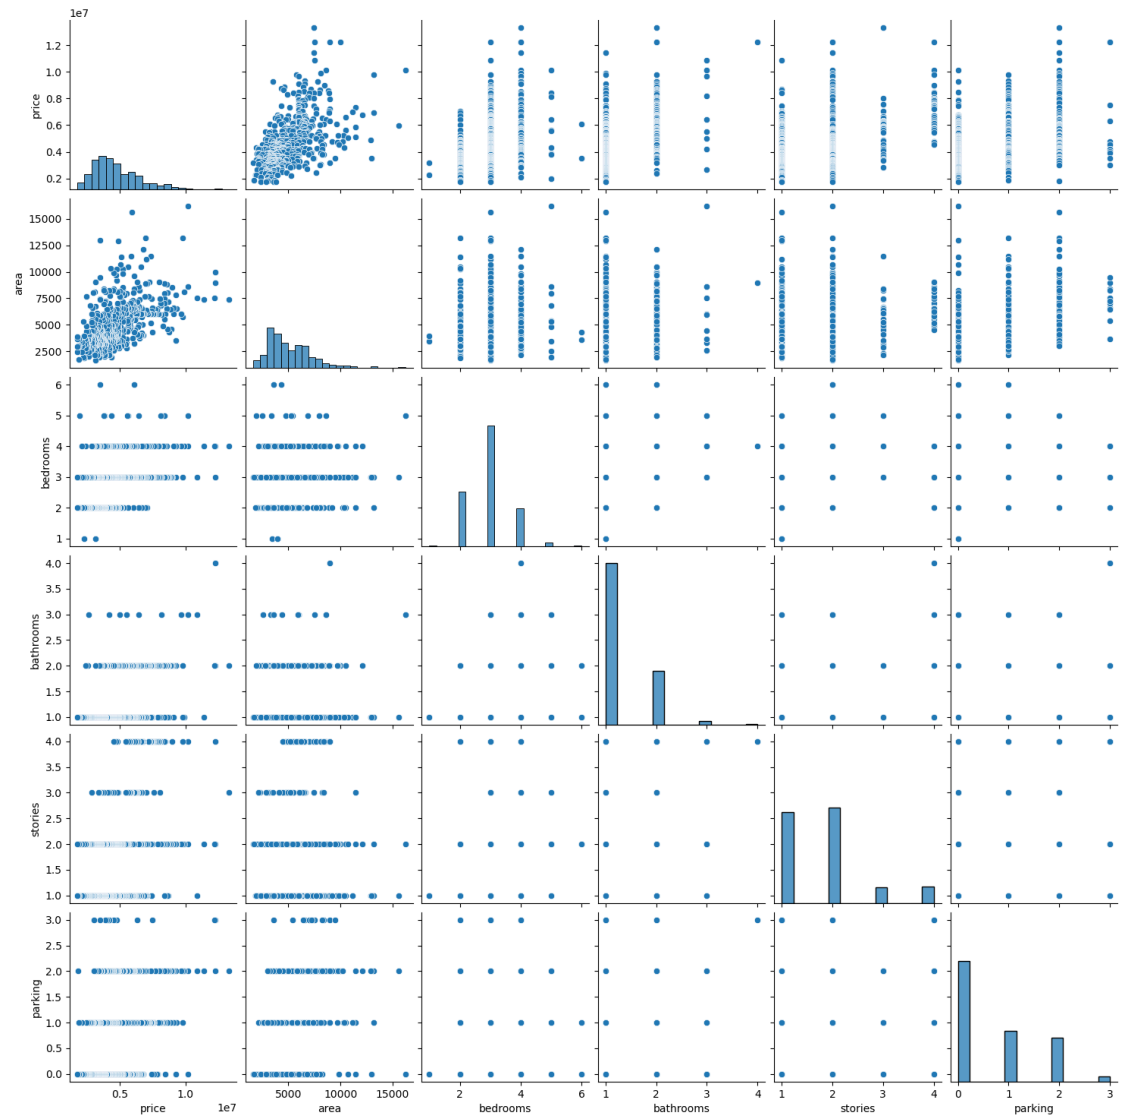
```
Out[14]: price          0  
         area          0  
         bedrooms      0  
         bathrooms     0  
         stories       0  
         mainroad       0  
         guestroom      0  
         basement       0  
         hotwaterheating 0  
         airconditioning 0  
         parking        0  
         prefarea       0  
         furnishingstatus 0  
         dtype: int64
```

In [15]: `## VISUALIZATION ##`

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.pairplot(df)
```

Out[15]: `<seaborn.axisgrid.PairGrid at 0x1955c866160>`



In [16]:  `## CLASSIFICATION OF CATEGORICAL DATA ##`

```
print(df.mainroad.value_counts())
print(df.guestroom.value_counts())
print(df.basement.value_counts())
print(df.hotwaterheating.value_counts())
print(df.airconditioning.value_counts())
print(df.prefarea.value_counts())
print(df.furnishingstatus.value_counts())
```

```
yes    468
no      77
Name: mainroad, dtype: int64
no     448
yes     97
Name: guestroom, dtype: int64
no     354
yes    191
Name: basement, dtype: int64
no     520
yes     25
Name: hotwaterheating, dtype: int64
no     373
yes    172
Name: airconditioning, dtype: int64
no     417
yes    128
Name: prefarea, dtype: int64
semi-furnished    227
unfurnished       178
furnished         140
Name: furnishingstatus, dtype: int64
```

In [17]:  `df['bedrooms'].value_counts()`

```
Out[17]: 3    300
         2    136
         4     95
         5     10
         6      2
         1      2
         Name: bedrooms, dtype: int64
```

In [18]:  `df['bathrooms'].value_counts()`

```
Out[18]: 1    401
         2    133
         3     10
         4      1
         Name: bathrooms, dtype: int64
```

```
In [19]: df['stories'].value_counts()
```

```
Out[19]: 2    238
         1    227
         4     41
         3     39
         Name: stories, dtype: int64
```

```
In [20]: df['parking'].value_counts()
```

```
Out[20]: 0    299
         1    126
         2    108
         3     12
         Name: parking, dtype: int64
```

```
In [21]: df['area'].value_counts()
```

```
Out[21]: 6000    24
         3000    14
         4500    13
         4000    11
         5500     9
         ..
         6862     1
         4815     1
         9166     1
         6321     1
         3620     1
         Name: area, Length: 284, dtype: int64
```

```
In [23]: ## CHANGING THE CATEGORICAL ATTRIBUTES INTO NUMERIC DATA FOR BETTER ANALYSIS
```

```
df.replace({'mainroad':{'yes':0,'no':1}},inplace=True)
df.replace({'guestroom':{'yes':0,'no':1}},inplace=True)
df.replace({'basement':{'yes':0,'no':1}},inplace=True)
df.replace({'hotwaterheating':{'yes':0,'no':1}},inplace=True)
df.replace({'airconditioning':{'yes':0,'no':1}},inplace=True)
df.replace({'prefarea.value':{'yes':0,'no':1}},inplace=True)
df.replace({'furnishingstatus':{'furnished':0,'semi-furnished':1,'unfurnished':2}},inplace=True)
```

In [24]:

```
## hot encoding ##
df = pd.get_dummies(df,drop_first=True)
df
```

Out[24]:

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hot
0	13300000	7420	4	2	3	0	1	1	
1	12250000	8960	4	4	4	0	1	1	
2	12250000	9960	3	2	2	0	1	0	
3	12215000	7500	4	2	2	0	1	0	
4	11410000	7420	4	1	2	0	0	0	
...
540	1820000	3000	2	1	1	0	1	0	
541	1767150	2400	3	1	1	1	1	1	
542	1750000	3620	2	1	1	0	1	1	
543	1750000	2910	3	1	1	1	1	1	
544	1750000	3850	3	1	2	0	1	1	

545 rows × 13 columns

```
In [25]: ▶ ## Splitting the data ##  
  
x = df.drop(['price'], axis=1)  
y = df['price']  
print(len(x), len(y))  
print(x)  
print(y)  
print(x.shape)  
print(y.shape)
```



```

545 545
      area  bedrooms  bathrooms  stories  mainroad  guestroom  basement
\
0    7420         4         2         3         0         1         1
1    8960         4         4         4         0         1         1
2    9960         3         2         2         0         1         0
3    7500         4         2         2         0         1         0
4    7420         4         1         2         0         0         0
..    ...         ...         ...         ...         ...         ...         ...
540   3000         2         1         1         0         1         0
541   2400         3         1         1         1         1         1
542   3620         2         1         1         0         1         1
543   2910         3         1         1         1         1         1
544   3850         3         1         2         0         1         1

      hotwaterheating  airconditioning  parking  furnishingstatus  prefa
ea_yes
0         1         0         2         0
1
1         1         0         3         0
0
2         1         1         2         1
1
3         1         0         3         0
1
4         1         0         2         0
0
..         ...         ...         ...         ...
...
540         1         1         2         2
0
541         1         1         0         1
0
542         1         1         0         2
0
543         1         1         0         0
0
544         1         1         0         2
0

```

```
[545 rows x 12 columns]
```

```

0    13300000
1    12250000
2    12250000
3    12215000
4    11410000

```

```

...
540    1820000
541    1767150
542    1750000
543    1750000
544    1750000

```

```

Name: price, Length: 545, dtype: int64
(545, 12)
(545,)

```

```
In [26]: ## Training and Test Data ##  
from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=1/10,rand
```

```
In [27]: ## Linear Regression ##  
  
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(x_train,y_train)  
model.score(x_test,y_test)
```

Out[27]: 0.6167018298753526

```
In [28]: from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()  
regressor.fit(x_train, y_train)
```


Out[28]: LinearRegression()

```
In [29]: t_data_predic = regressor.predict(x_train)
```

```
In [30]: ## Error Calculation ##  
  
from sklearn import metrics  
error_score = metrics.r2_score(y_train, t_data_predic)  
print("R squared Error : ", error_score)
```


R squared Error : 0.6836517746805287

```
In [31]: ► ## Plotting THE data ##  
  
plt.scatter(y_train, t_data_predic)  
plt.xlabel("Actual Price of Car")  
plt.ylabel("Predicted Price of Car")  
plt.title(" Actual Prices vs Predicted Prices")  
plt.show()
```




```
In [32]: ► ## prediction on Test data ##  
t_data_predic = regressor.predict(x_test)
```

```
In [33]: ▶ plt.scatter(y_test, t_data_predic)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")
plt.show()
```



```
In [35]: ▶ plt.scatter(y_test, t_data_predic)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title(" Actual Prices vs Predicted Prices")
plt.plot([min(y_train), max(y_train)], [min(y_train), max(y_train)], 'r--')
plt.show()
```



```
In [36]: ▶ b = regressor.coef_
print("Coefficient :",b)
```

```
Coefficient : [ 2.47133589e+02  1.19610191e+05  1.01501859e+06  4.425545
54e+05
-4.13328625e+05 -4.03165037e+05 -3.37574049e+05 -1.05647054e+06
-8.85606823e+05  2.63424448e+05 -1.96010630e+05  6.69004383e+05]
```

```
In [37]: ▶ a = regressor.intercept_
print("Intercept :",a)
```

```
Intercept : 3132124.8084780285
```